


Accepted Manuscript

Gaussian process classification for prediction of in-hospital mortality among preterm infants

Olli-Pekka Rinta-Koski, Simo Särkkä, Jaakko Hollmén, Markus Leskinen, Sture Andersson

PII: S0925-2312(18)30208-X
DOI: [10.1016/j.neucom.2017.12.064](https://doi.org/10.1016/j.neucom.2017.12.064)
Reference: NEUCOM 19352



To appear in: *Neurocomputing*

Received date: 4 July 2017
Revised date: 10 November 2017
Accepted date: 21 December 2017

Please cite this article as: Olli-Pekka Rinta-Koski, Simo Särkkä, Jaakko Hollmén, Markus Leskinen, Sture Andersson, Gaussian process classification for prediction of in-hospital mortality among preterm infants, *Neurocomputing* (2018), doi: [10.1016/j.neucom.2017.12.064](https://doi.org/10.1016/j.neucom.2017.12.064)

This is a PDF file of an unedited manuscript that has been accepted for publication. As a service to our customers we are providing this early version of the manuscript. The manuscript will undergo copyediting, typesetting, and review of the resulting proof before it is published in its final form. Please note that during the production process errors may be discovered which could affect the content, and all legal disclaimers that apply to the journal pertain.

Gaussian process classification for prediction of in-hospital mortality among preterm infants

Olli-Pekka Rinta-Koski^{a,*}, Simo Särkkä^b, Jaakko Hollmén^a, Markus Leskinen^c,
Sture Andersson^c

^a*Aalto University, Department of Computer Science, PO Box 15400, FI-00076 Aalto, Finland*

^b*Aalto University, Department of Electrical Engineering and Automation, PO Box 12200, FI-00076 Aalto, Finland*

^c*University of Helsinki, and Helsinki University Hospital, PO Box 140, FI-00029 HUS, Finland*

Abstract

We present a method for predicting preterm infant in-hospital mortality using Bayesian Gaussian process classification. We combined features extracted from sensor measurements, made during the first 72 hours of care for 598 Very Low Birth Weight infants of birth weight <1500 g, with standard clinical features calculated on arrival at the Neonatal Intensive Care Unit. Time periods of 12, 18, 24, 36, 48, and 72 hours were evaluated. We achieved a classification result with area under the receiver operating characteristic curve of 0.948, which is in excess of the results achieved by using the clinical standard SNAP-II and SNAPPE-II scores.

Keywords: time series prediction; Gaussian process classification; very low birth weight infants; neonatal intensive care

1. Introduction

This article is related to the use of data-driven methods in the context of digital healthcare and health informatics [1, 2]. In particular, our aim is to develop

*Corresponding author

Email addresses: `olli-pekka.rinta-koski@aalto.fi` (Olli-Pekka Rinta-Koski), `simo.sarkka@aalto.fi` (Simo Särkkä), `jaakko.hollmen@aalto.fi` (Jaakko Hollmén), `markus.leskinen@hus.fi` (Markus Leskinen), `sture.andersson@hus.fi` (Sture Andersson)

machine learning methodology for integration of heterogeneous data sources in order to more accurately predict the survival chances of preterm infants during treatment in the Neonatal Intensive Care Unit (NICU). First, we combine the conventional scoring system used in clinical practice with data-driven prediction from raw sensor data. Second, we study the prediction accuracy when the clinical scores are completely replaced with measurement data. The development of new methods for predicting neonatal in-hospital mortality is important, because while the global under-five mortality rate has dropped 53% since 1990, the proportion of neonatal deaths is projected to increase from 45% in 2015 to 52% by 2030 [3]. The incidence of certain complications (e.g. necrotizing enterocolitis) increases with the survival of preterm infants who previously would have died before the onset of these problems, emphasizing the need for developing new methods and strategies for neonatal intensive care [4]. Furthermore, data-only prediction is extremely important in clinical work, because the determination of the conventional scores is labor-intensive and requires that a specific set of diagnostic markers is available.

Routinely available markers of risk – sex, birth weight, and gestational age – fail to predict observed variation of mortality in NICUs [5]. This has prompted development of illness severity scores, such as SNAP-II and SNAPPE-II [6], which add laboratory results and physiological measurements of vital signs to perinatal risk factors in order to better predict morbidity and mortality. These risk scores were developed when patient records were mostly collected by hand, relying on simplified presentation of physiological data such as lowest temperature and mean blood pressure. Current patient information systems and patient monitors have made collection of detailed medical data much easier. We hypothesized that time series data of vital signs would help to identify patients at risk and, when combined with traditional risk scores, would result in increased predictive power.

The machine learning methodology that we use is based on the use of Gaussian process (GP) classification [7] with features extracted from raw cardiac, arterial and oximeter sensor measurements in addition to the clinical scores,

35 gestational age at birth, and birth weight. Our motivation for studying GP
 classifiers in this context stems from two properties of GPs. First, they are gen-
 uine probabilistic models [7] and can provide information on how certain we are
 about the answer. This feature is inherent in GPs whereas, for example, for sup-
 port vector machines (SVMs) [8] the uncertainty needs to be estimated with an
 40 additional model on the basic SVM [9]. Second, an even more important prop-
 erty is that GPs can flexibly be combined with first principles models [10, 11].
 The resulting latent force models (LFMs) have a huge potential in medical appli-
 cations especially due to their connection with time-series models used in sensor
 signal processing [12, 13, 14]. As shown in these papers, it is even possible to
 45 see that GPs models are solutions to certain stochastic partial differential equa-
 tions, which not only allow for the combination with first-principles physical
 models, but also enable the use of Kalman filtering and other Bayesian filtering
 methods [15] for computationally efficient implementation of GP classifiers. GP
 classifiers have been previously used in health data analysis in (adult) Intensive
 50 Care Units (ICU) [16, 17, 18] and machine learning methods have been applied
 to NICU data [19, 20].

The contribution of this paper is that using cross-validation we show that
 augmenting the staff-determined SNAP-II and SNAPPE-II scores with sensor
 measurements improves prediction accuracy over standard clinical measures.
 55 We also show that a data-driven prediction from measurements alone can lead
 to better prediction accuracy than SNAP-II and SNAPPE-II. The proposed
 approach gives the area under the receiver operating characteristic curve (AUC)
 0.946 for mortality prediction, which compares favourably with AUC 0.9151
 reported for logistic regression by Saria et al. [20], and AUC 0.913 for CRIB-
 60 II and AUC 0.907 for SNAPPE-II reported by Reid et al. [21]. Although it
 has previously been shown [6] that in-hospital mortality of preterm infants is
 strongly correlated with birth weight and gestational age at birth, we show that
 the prediction result achieved by using these two variables alone (Table 2) can
 be improved by adding features extracted from measurement time series.

65 This article is an extended version of the conference article “Prediction of

preterm infant mortality with Gaussian process classification” [22] presented at the 25th European Symposium on Artificial Neural Networks, Computational Intelligence and Machine Learning (ESANN 2017), in which we looked at data from the first postnatal 24 hours. Here the analysis has been extended to six
 70 different time periods ranging between the first 12 and 72 postnatal hours, three different kernels have been used with the GP classifier, and the classification performance has been compared to other classifiers.

2. Materials and methods

2.1. NICU database

75 The NICU at Helsinki University Hospital has been collecting patient data in a database since 1999. Data include measurements of clinical parameters such as oxygen saturation by pulse oximetry (SpO_2) and supplemental oxygen levels, observations made by staff, and clinical outcomes. Our study cohort includes 2059 Very Low Birth Weight (VLBW) infants (birth weight <1500 g)
 80 admitted between 1999–2013. Median gestational age at birth was 202 days (H28+6 weeks) and median birth weight was 1102 g.

The NICU database contains data recorded from equipment interfaces, as well as notations made by hand. Automatically gathered data consists of 111 different variables taken from monitor outputs of equipment used in the NICU.
 85 As the monitoring equipment and clinical guidelines have varied during the 15 year period under which the data has been stored, not all data is available for all 2059 patients.

2.2. Preprocessing and feature extraction

90 For the experiment, we decided to study the first 72 hours from delivery to see whether the time series data gathered during that period has predictive power. Most in-hospital deaths occur within the first week; median in this dataset is 5 days. There are 598 patients in the dataset for whom there is complete data from the first 72 hours of their NICU stay for each of these seven

variables: gestational age at birth, birth weight, systolic, mean, and diastolic
 95 arterial blood pressure, heart rate measured by electrocardiography (ECG), and
 SpO₂. If for some sensor signal there were only a few measurements available,
 the patient data was considered incomplete. Patients that died before the end
 of 72 hour period were excluded as well. The in-hospital mortality rate of this
 subset is 9% (53 patients), which is also the mortality rate in the full cohort.
 100 In addition to the full 72 hour period, we also looked at the first 12, 18, 24, 36,
 and 48 hour periods.

The data was preprocessed by removing out-of-range values caused by, for
 example, misplaced or removed sensors and monitoring equipment drifting out
 of calibration from the time series.

105 For feature extraction, mean and standard deviation were calculated from
 each of the following time series for each patient: systolic, mean, and arte-
 rial blood pressure, ECG heart rate, and SpO₂. SNAP-II score, SNAPPE-II
 score, gestational age at birth, and birth weight were directly used as features.
 We chose not to use any more complicated features such as signal derivatives,
 110 because the signals streams were very sparse and noisy, and reliably estimat-
 ing the signal derivatives would have required us to use Kalman filter type
 of methods [15], which we wanted to avoid at this stage in order to keep the
 preprocessing simple and robust.

2.3. Gaussian process classifier

115 We used a GP [7] classifier with a probit measurement model:

$$f(\mathbf{x}) \sim \mathcal{GP}(0, k(\mathbf{x}, \mathbf{x}')) \quad p(y_i | f(\mathbf{x}_i)) = \int_{-\infty}^{y_i f(\mathbf{x}_i)} \mathcal{N}(z | 0, 1) dz \quad (1)$$

where the classes are labeled as $y_i \in \{-1, 1\}$. This choice of the measurement
 model is standard in GP literature [7] and is supported by most GP software
 packages such as the GPstuff Toolbox [23].

The kernel was a sum of squared exponential (or radial basis function) kernel,
 120 linear kernel, and constant kernel:

$$k(\mathbf{x}, \mathbf{x}') = \sigma_{\text{se}}^2 \exp\left(-\frac{1}{2}(\mathbf{x} - \mathbf{x}')^\top \Lambda^{-1}(\mathbf{x} - \mathbf{x}')\right) + \mathbf{x}^\top \Sigma \mathbf{x}' + \sigma^2, \quad (2)$$

where $\Lambda = \text{diag}(l_1^2, \dots, l_d^2)$ and $\Sigma = \text{diag}(\sigma_1^2, \dots, \sigma_d^2)$. The rationale behind this kernel choice is that the constant and the linear parts of the kernel aim at capturing the bias and the linear trend in the problem, respectively. In order to capture the non-linear effects, we add the squared exponential kernel with
125 the automatic relevance determination prior, which is a commonly used general covariance function in Gaussian process regression [7]. This kind of 3-part covariance functions have also been recently used in medical applications [24, 25].

For comparison purposes, we also used the Matérn kernel with $\nu = 3/2$ (M32) and $\nu = 5/2$ (M52) [7, 23, 24, 25] replacing the squared exponential
130 kernel in the 3-part kernel:

$$k_{\nu=3/2}(\mathbf{x}, \mathbf{x}') = \sigma_m^2 (1 + \sqrt{3}\|\mathbf{x} - \mathbf{x}'\|) \exp(-\sqrt{3}\|\mathbf{x} - \mathbf{x}'\|) \quad (3)$$

$$k_{\nu=5/2}(\mathbf{x}, \mathbf{x}') = \sigma_m^2 (1 + \sqrt{5}\|\mathbf{x} - \mathbf{x}'\| + \frac{5\|\mathbf{x} - \mathbf{x}'\|^2}{3}) \exp(-\sqrt{5}\|\mathbf{x} - \mathbf{x}'\|) \quad (4)$$

For training the classifier we used the GPstuff Toolbox [23] with Laplace approximation on the latent variables and circular composite design (CCD) integration over the hyperparameters. The CCD method has advantages over, for
135 example, marginal likelihood maximization due to better handling of uncertainty. In particular, the method approximates the integration over the hyperparameters instead of using a plug-in point-estimate, which ensures that the uncertainty is computed in a proper Bayesian way.

In order to evaluate the performance of the classifiers we used stratified 8-fold
140 cross-validation (CV) which takes the class priors into account when forming the partitions. Cross-validation was used to estimate the classification accuracy, precision, specificity, and sensitivity as well as receiver operating characteristic (ROC) curve [26] and the area under the ROC curve (AUC) [27]. In order to reduce the variance of CV, we repeated each CV run 8 times and averaged the
145 results.

2.4. Comparison with other classifiers

We used the following classifiers in comparison with the GP classifier:

- *SNAP-II/SNAPPE-II thresholding.* Thresholding using only the SNAP-II or SNAPPE-II scores (one at a time) was used to classify the patients. The class boundary was set using one of two rules. In the first case, the maximum accuracy achieved with the training set was used to set the class boundary. In the second case, the maximum value of the Youden index [28] was used. This gave us four different rule-score combinations.
- *Support vector machine classifier.* A linear SVM classifier [8] was used as the classifier and the posterior probability estimates were obtained with Platt scaling [9]. The ROC curve was calculated by sweeping the class boundary from 0 to 1. The prediction was given by setting the class boundary to 0.5.
- *Linear probit model.* A linear model with a probit link function was implemented by using a constant plus a linear kernel in a GP classification model. The model was trained using the GPstuff Toolbox. The integration over the hyperparameters was performed using the CCD method [29].
- *Random classifier.* This classifier assigns the class at random weighted by class prior probabilities of the training set.
- *Majority classifier.* This classifier simply assumes that all patients belong to the larger (survivor) class.

3. Results

3.1. Classification with SNAP-II and SNAPPE-II scores

First, we tested the performance of the classifiers using only SNAP-II and SNAPPE-II scores with gestational age at birth and birth weight. Although this information is equal to what the scores are traditionally computed from, as can be seen in Tables 1 and 2, the GP classifier is able to achieve a better AUC (0.933) than the clinical standard SNAP-II (AUC 0.860) and SNAPPE-II (AUC 0.878) scores, with all variants (sum, M32, M52) giving practically the same result and linear probit classifier at a just slightly lower AUC (0.921).

Table 1: Reference results.

	Acc	PPV	Sens	Spec	AUC
SNAPPE A	0.914 (0.00)	0.898 (0.05)	0.056 (0.02)	0.998 (0.00)	0.875 (0.00)
SNAPPE Y	0.737 (0.01)	0.248 (0.01)	0.923 (0.02)	0.719 (0.01)	0.875 (0.00)
SNAP A	0.909 (0.00)	0.677 (0.08)	0.062 (0.02)	0.993 (0.00)	0.859 (0.00)
SNAP Y	0.713 (0.01)	0.227 (0.01)	0.895 (0.02)	0.695 (0.01)	0.859 (0.00)
Random	0.839 (0.01)	0.091 (0.02)	0.091 (0.02)	0.913 (0.01)	0.500 (0.00)
Majority	0.910 (0.00)	1.000 (0.00)	0.000 (0.00)	1.000 (0.00)	0.500 (0.00)

SNAP/SNAPPE = SNAP-II/SNAPPE-II with optimal (cross-validated) thresholding (A = maximal accuracy, Y = Youden index), Majority = trivial classifier that assumes all patients survive, Random = class picked at random weighted by training set class priors. Acc = accuracy, PPV = positive predictive value, Sens = sensitivity, Spec = specificity, AUC = area under the receiver operating characteristic curve. Values in parentheses indicate the associated standard error. Results in all tables in descending order by AUC.

Table 2: Results using only SNAP-II, SNAPPE-II, gestational age at birth, and birth weight.

	Acc	PPV	Sens	Spec	AUC
GP M32	0.918 (0.00)	0.611 (0.04)	0.360 (0.03)	0.974 (0.00)	0.933 (0.00)
GP	0.919 (0.00)	0.618 (0.05)	0.351 (0.03)	0.975 (0.00)	0.933 (0.00)
GP M52	0.919 (0.00)	0.615 (0.04)	0.358 (0.03)	0.974 (0.00)	0.933 (0.00)
Linear	0.914 (0.00)	0.579 (0.05)	0.260 (0.03)	0.978 (0.00)	0.921 (0.00)
SVM	0.907 (0.00)	0.857 (0.06)	0.020 (0.01)	0.994 (0.00)	0.644 (0.01)

Next, we used all available signals with the GP classifier in order to get an upper bound on the achievable performance. All the available features were used as classifier inputs, in other words, SNAP-II, SNAPPE-II, gestational age at birth, birth weight, and the mean and standard deviation of each of the following: systolic, mean, and diastolic arterial blood pressure, ECG heart rate, and SpO₂.

Table 3 (all available features) shows GP prediction results using all available features with three different kernels (sum, M32, and M52). Kernel choice had a negligible effect. The highest AUC (0.948) was achieved with 48h data and the sum kernel. All AUC values from predictions with all three kernels for time periods between 36h and 72h were within 0.007. Shortening the range of time

series data has a slight negative effect on the AUC values, with 12h data and the sum kernel yielding AUC 0.924. However, as the range decreases, there is a drop in both positive predictive value (PPV), from 0.708 to 0.598, and sensitivity, 190 from 0.463 to 0.283. SVM and the linear probit model give similar results to GP (Tables 4 and 5) in many of the cases, but with shorter ranges the GP models give slightly better results.

Table 3: GP prediction results using all available features and three different kernels (sum, M32, M52).

	Acc	PPV	Sens	Spec	AUC
48h GP	0.930 (0.00)	0.660 (0.03)	0.463 (0.02)	0.975 (0.00)	0.948 (0.00)
48h GPm32	0.928 (0.00)	0.649 (0.03)	0.445 (0.02)	0.975 (0.00)	0.947 (0.00)
48h GPm52	0.928 (0.00)	0.657 (0.03)	0.442 (0.02)	0.976 (0.00)	0.946 (0.00)
36h GPm32	0.925 (0.00)	0.667 (0.03)	0.391 (0.02)	0.977 (0.00)	0.945 (0.00)
72h GP	0.932 (0.00)	0.708 (0.03)	0.449 (0.02)	0.980 (0.00)	0.942 (0.01)
72h GPm52	0.933 (0.00)	0.717 (0.03)	0.453 (0.02)	0.980 (0.00)	0.942 (0.01)
36h GPm52	0.925 (0.00)	0.669 (0.03)	0.390 (0.02)	0.977 (0.00)	0.942 (0.01)
72h GPm32	0.934 (0.00)	0.727 (0.03)	0.452 (0.02)	0.981 (0.00)	0.942 (0.01)
36h GP	0.924 (0.00)	0.670 (0.03)	0.389 (0.02)	0.977 (0.00)	0.941 (0.01)
24h GPm32	0.918 (0.00)	0.591 (0.04)	0.332 (0.03)	0.976 (0.00)	0.931 (0.01)
24h GPm52	0.919 (0.00)	0.596 (0.04)	0.331 (0.03)	0.977 (0.00)	0.930 (0.01)
18h GPm32	0.920 (0.00)	0.671 (0.03)	0.312 (0.02)	0.981 (0.00)	0.929 (0.01)
24h GP	0.919 (0.00)	0.624 (0.03)	0.335 (0.02)	0.976 (0.00)	0.929 (0.00)
18h GP	0.920 (0.00)	0.655 (0.04)	0.328 (0.02)	0.979 (0.00)	0.928 (0.01)
18h GPm52	0.922 (0.00)	0.682 (0.03)	0.335 (0.02)	0.980 (0.00)	0.928 (0.01)
12h GP	0.915 (0.00)	0.598 (0.04)	0.283 (0.02)	0.977 (0.00)	0.924 (0.01)
12h GPm32	0.913 (0.00)	0.590 (0.04)	0.295 (0.02)	0.974 (0.00)	0.923 (0.01)
12h GPm52	0.913 (0.00)	0.581 (0.03)	0.297 (0.02)	0.974 (0.00)	0.921 (0.01)

Finally, Table 6 shows the results for all non-reference classifiers using all available features. GP kernel choice had negligible effect. The linear probit 195 model performs worse than GP with 12h and 18h data. With longer time series, GP and the linear probit model have roughly equal performance. Even the lowest AUC (0.901), given by the SVM classifier with 12h data, is better than SNAP-II/SNAPPE-II thresholding (AUC 0.859...0.875).

Table 4: SVM prediction results using all available features.

	Acc	PPV	Sens	Spec	AUC
36h	0.931 (0.00)	0.708 (0.03)	0.443 (0.02)	0.979 (0.00)	0.947 (0.00)
48h	0.931 (0.00)	0.725 (0.03)	0.433 (0.02)	0.980 (0.00)	0.943 (0.00)
24h	0.923 (0.00)	0.665 (0.03)	0.340 (0.02)	0.981 (0.00)	0.930 (0.01)
72h	0.931 (0.00)	0.746 (0.03)	0.384 (0.02)	0.984 (0.00)	0.924 (0.01)
18h	0.918 (0.00)	0.692 (0.04)	0.259 (0.02)	0.984 (0.00)	0.922 (0.00)
12h	0.910 (0.00)	0.564 (0.04)	0.182 (0.02)	0.982 (0.00)	0.901 (0.01)

Table 5: Linear model prediction results using all available features.

	Acc	PPV	Sens	Spec	AUC
48h	0.926 (0.00)	0.645 (0.02)	0.463 (0.02)	0.971 (0.00)	0.949 (0.00)
36h	0.924 (0.00)	0.652 (0.03)	0.402 (0.02)	0.975 (0.00)	0.949 (0.00)
72h	0.934 (0.00)	0.720 (0.03)	0.475 (0.02)	0.979 (0.00)	0.944 (0.00)
24h	0.919 (0.00)	0.610 (0.04)	0.348 (0.02)	0.976 (0.00)	0.931 (0.01)
18h	0.922 (0.00)	0.675 (0.04)	0.326 (0.02)	0.981 (0.00)	0.927 (0.01)
12h	0.913 (0.00)	0.583 (0.04)	0.279 (0.02)	0.976 (0.00)	0.916 (0.01)

3.2. Classification with reduced feature sets

200 To find out how the classifiers perform with reduced feature sets, we tested the classifiers without SNAP-II and SNAPPE-II scores (Table 7) and finally with sensor signals only (dropping also gestational age at birth and birth weight, Table 8).

Without SNAP-II/SNAPPE-II, the linear probit model and GP perform 205 equally well with time periods of at least 36h (AUC 0.943...0.946). All classifiers outperform the reference results (Table 1), with the exception of SVM with 12h data (AUC 0.874) which achieves a result comparable with SNAPPE-II thresholding.

210 Table 8 shows prediction results using only time series data. The best classifier is GP (all kernels) with 48h data (AUC 0.925...0.926) but linear classifiers with 48h and 72h data as well as GP with 72h data perform almost equally well. Whereas the GP kernel choice is again practically immaterial, both AUC and sensitivity increase as the time series grows longer, AUC from 0.787 (12h data, M32 kernel) to 0.926 (48h data, sum kernel). Interestingly, 72h data gives 215 slightly lower AUC values than 48h data (AUC 0.915...0.919 vs. 0.925...0.926),

Table 6: Prediction results using all available features. Comparison of GP, SVM, and linear probit model.

	Acc	PPV	Sens	Spec	AUC
48h Linear	0.926 (0.00)	0.645 (0.02)	0.463 (0.02)	0.971 (0.00)	0.949 (0.00)
36h Linear	0.924 (0.00)	0.652 (0.03)	0.402 (0.02)	0.975 (0.00)	0.949 (0.00)
48h GP	0.930 (0.00)	0.660 (0.03)	0.463 (0.02)	0.975 (0.00)	0.948 (0.00)
36h SVM	0.931 (0.00)	0.708 (0.03)	0.443 (0.02)	0.979 (0.00)	0.947 (0.00)
48h GPm32	0.928 (0.00)	0.649 (0.03)	0.445 (0.02)	0.975 (0.00)	0.947 (0.00)
48h GPm52	0.928 (0.00)	0.657 (0.03)	0.442 (0.02)	0.976 (0.00)	0.946 (0.00)
36h GPm32	0.925 (0.00)	0.667 (0.03)	0.391 (0.02)	0.977 (0.00)	0.945 (0.00)
72h Linear	0.934 (0.00)	0.720 (0.03)	0.475 (0.02)	0.979 (0.00)	0.944 (0.00)
48h SVM	0.931 (0.00)	0.725 (0.03)	0.433 (0.02)	0.980 (0.00)	0.943 (0.00)
72h GP	0.932 (0.00)	0.708 (0.03)	0.449 (0.02)	0.980 (0.00)	0.942 (0.01)
72h GPm52	0.933 (0.00)	0.717 (0.03)	0.453 (0.02)	0.980 (0.00)	0.942 (0.01)
36h GPm52	0.925 (0.00)	0.669 (0.03)	0.390 (0.02)	0.977 (0.00)	0.942 (0.01)
72h GPm32	0.934 (0.00)	0.727 (0.03)	0.452 (0.02)	0.981 (0.00)	0.942 (0.01)
36h GP	0.924 (0.00)	0.670 (0.03)	0.389 (0.02)	0.977 (0.00)	0.941 (0.01)
24h GPm32	0.918 (0.00)	0.591 (0.04)	0.332 (0.03)	0.976 (0.00)	0.931 (0.01)
24h Linear	0.919 (0.00)	0.610 (0.04)	0.348 (0.02)	0.976 (0.00)	0.931 (0.01)
24h SVM	0.923 (0.00)	0.665 (0.03)	0.340 (0.02)	0.981 (0.00)	0.930 (0.01)
24h GPm52	0.919 (0.00)	0.596 (0.04)	0.331 (0.03)	0.977 (0.00)	0.930 (0.01)
18h GPm32	0.920 (0.00)	0.671 (0.03)	0.312 (0.02)	0.981 (0.00)	0.929 (0.01)
24h GP	0.919 (0.00)	0.624 (0.03)	0.335 (0.02)	0.976 (0.00)	0.929 (0.00)
18h GP	0.920 (0.00)	0.655 (0.04)	0.328 (0.02)	0.979 (0.00)	0.928 (0.01)
18h GPm52	0.922 (0.00)	0.682 (0.03)	0.335 (0.02)	0.980 (0.00)	0.928 (0.01)
18h Linear	0.922 (0.00)	0.675 (0.04)	0.326 (0.02)	0.981 (0.00)	0.927 (0.01)
72h SVM	0.931 (0.00)	0.746 (0.03)	0.384 (0.02)	0.984 (0.00)	0.924 (0.01)
12h GP	0.915 (0.00)	0.598 (0.04)	0.283 (0.02)	0.977 (0.00)	0.924 (0.01)
12h GPm32	0.913 (0.00)	0.590 (0.04)	0.295 (0.02)	0.974 (0.00)	0.923 (0.01)
18h SVM	0.918 (0.00)	0.692 (0.04)	0.259 (0.02)	0.984 (0.00)	0.922 (0.00)
12h GPm52	0.913 (0.00)	0.581 (0.03)	0.297 (0.02)	0.974 (0.00)	0.921 (0.01)
12h Linear	0.913 (0.00)	0.583 (0.04)	0.279 (0.02)	0.976 (0.00)	0.916 (0.01)
12h SVM	0.910 (0.00)	0.564 (0.04)	0.182 (0.02)	0.982 (0.00)	0.901 (0.01)

but with better PPV and sensitivity (PPV 0.804...0.813 vs. 0.639...0.645, sensitivity 0.347...0.361 vs. 0.320...0.335). Time periods shorter than 36h do not give better than reference results with any of the classifiers.

The best SVM result (AUC 0.899, 48h data) equals the performance of GP and linear classifiers with 36h data, but loses to both with time periods of 48h and 72h. SVM performance degrades markedly with 24h and shorter data, not

beating even the reference (SNAP-II/SNAPPE-II) classifiers.

Table 7: Prediction results using all available features except SNAP-II and SNAPPE-II.

	Acc	PPV	Sens	Spec	AUC
48h GPm32	0.928 (0.00)	0.680 (0.03)	0.442 (0.02)	0.976 (0.00)	0.947 (0.00)
48h Linear	0.926 (0.00)	0.655 (0.03)	0.464 (0.02)	0.971 (0.00)	0.947 (0.00)
48h GP	0.929 (0.00)	0.678 (0.03)	0.445 (0.02)	0.976 (0.00)	0.946 (0.00)
48h GPm52	0.926 (0.00)	0.668 (0.03)	0.433 (0.02)	0.975 (0.00)	0.946 (0.00)
72h GPm32	0.934 (0.00)	0.735 (0.02)	0.445 (0.02)	0.981 (0.00)	0.946 (0.00)
36h GP	0.922 (0.00)	0.634 (0.03)	0.400 (0.02)	0.973 (0.00)	0.945 (0.00)
36h GPm52	0.922 (0.00)	0.617 (0.03)	0.403 (0.02)	0.973 (0.00)	0.945 (0.00)
36h GPm32	0.923 (0.00)	0.627 (0.03)	0.412 (0.02)	0.973 (0.00)	0.945 (0.00)
72h Linear	0.934 (0.00)	0.705 (0.02)	0.503 (0.02)	0.976 (0.00)	0.945 (0.00)
36h Linear	0.927 (0.00)	0.655 (0.03)	0.457 (0.02)	0.973 (0.00)	0.945 (0.00)
72h GPm52	0.933 (0.00)	0.723 (0.02)	0.448 (0.02)	0.980 (0.00)	0.944 (0.01)
72h GP	0.932 (0.00)	0.713 (0.03)	0.439 (0.02)	0.980 (0.00)	0.943 (0.01)
48h SVM	0.927 (0.00)	0.691 (0.03)	0.413 (0.02)	0.978 (0.00)	0.941 (0.00)
36h SVM	0.930 (0.00)	0.699 (0.03)	0.413 (0.02)	0.981 (0.00)	0.941 (0.00)
24h Linear	0.924 (0.00)	0.639 (0.03)	0.396 (0.03)	0.976 (0.00)	0.931 (0.00)
24h GP	0.921 (0.00)	0.608 (0.03)	0.357 (0.03)	0.977 (0.00)	0.930 (0.00)
24h GPm52	0.919 (0.00)	0.590 (0.03)	0.349 (0.03)	0.975 (0.00)	0.930 (0.00)
24h GPm32	0.918 (0.00)	0.570 (0.04)	0.343 (0.03)	0.975 (0.00)	0.930 (0.00)
18h GPm52	0.921 (0.00)	0.674 (0.03)	0.328 (0.02)	0.980 (0.00)	0.925 (0.01)
18h GPm32	0.921 (0.00)	0.670 (0.03)	0.333 (0.02)	0.980 (0.00)	0.925 (0.01)
18h GP	0.923 (0.00)	0.678 (0.03)	0.345 (0.02)	0.981 (0.00)	0.924 (0.01)
72h SVM	0.932 (0.00)	0.742 (0.03)	0.405 (0.02)	0.984 (0.00)	0.923 (0.01)
18h Linear	0.921 (0.00)	0.659 (0.03)	0.337 (0.02)	0.979 (0.00)	0.920 (0.01)
24h SVM	0.924 (0.00)	0.733 (0.03)	0.306 (0.02)	0.985 (0.00)	0.919 (0.01)
12h GPm32	0.914 (0.00)	0.586 (0.04)	0.258 (0.02)	0.979 (0.00)	0.914 (0.01)
12h GPm52	0.913 (0.00)	0.591 (0.04)	0.253 (0.02)	0.979 (0.00)	0.913 (0.01)
12h GP	0.912 (0.00)	0.562 (0.04)	0.253 (0.02)	0.977 (0.00)	0.912 (0.01)
18h SVM	0.921 (0.00)	0.738 (0.03)	0.259 (0.02)	0.987 (0.00)	0.907 (0.01)
12h Linear	0.915 (0.00)	0.594 (0.04)	0.256 (0.02)	0.981 (0.00)	0.900 (0.01)
12h SVM	0.909 (0.00)	0.622 (0.05)	0.084 (0.01)	0.990 (0.00)	0.874 (0.01)

3.3. The effect of varying input combinations and time series lengths

Figure 1 shows the ROC curves for classifier results using all features, without SNAP-II/SNAPPE-II, and time series data only for the two extremes (12h, 72h) of time series lengths. The classification performance is improved with longer time series in all three cases.

Table 8: Prediction results using only time series data.

	Acc	PPV	Sens	Spec	AUC
48h GP	0.923 (0.00)	0.640 (0.03)	0.335 (0.02)	0.981 (0.00)	0.926 (0.00)
48h GPm52	0.923 (0.00)	0.639 (0.04)	0.322 (0.02)	0.981 (0.00)	0.925 (0.00)
48h GPm32	0.923 (0.00)	0.645 (0.04)	0.320 (0.02)	0.982 (0.00)	0.925 (0.00)
72h GPm32	0.932 (0.00)	0.813 (0.03)	0.347 (0.02)	0.989 (0.00)	0.919 (0.00)
72h GPm52	0.933 (0.00)	0.820 (0.03)	0.361 (0.02)	0.989 (0.00)	0.917 (0.00)
48h Linear	0.922 (0.00)	0.674 (0.04)	0.315 (0.02)	0.982 (0.00)	0.917 (0.00)
72h GP	0.932 (0.00)	0.804 (0.03)	0.352 (0.02)	0.989 (0.00)	0.915 (0.00)
72h Linear	0.933 (0.00)	0.798 (0.03)	0.360 (0.02)	0.988 (0.00)	0.913 (0.00)
36h GPm52	0.922 (0.00)	0.715 (0.04)	0.261 (0.02)	0.986 (0.00)	0.902 (0.01)
36h GPm32	0.922 (0.00)	0.716 (0.04)	0.259 (0.02)	0.987 (0.00)	0.902 (0.01)
36h GP	0.923 (0.00)	0.734 (0.04)	0.266 (0.02)	0.988 (0.00)	0.900 (0.01)
48h SVM	0.922 (0.00)	0.747 (0.04)	0.261 (0.02)	0.986 (0.00)	0.899 (0.01)
36h Linear	0.923 (0.00)	0.740 (0.04)	0.272 (0.02)	0.987 (0.00)	0.898 (0.01)
72h SVM	0.932 (0.00)	0.863 (0.03)	0.287 (0.02)	0.994 (0.00)	0.892 (0.01)
36h SVM	0.925 (0.00)	0.874 (0.03)	0.213 (0.02)	0.995 (0.00)	0.881 (0.01)
24h GPm32	0.918 (0.00)	0.701 (0.05)	0.194 (0.02)	0.990 (0.00)	0.868 (0.01)
24h GPm52	0.917 (0.00)	0.690 (0.05)	0.194 (0.02)	0.989 (0.00)	0.864 (0.01)
24h GP	0.919 (0.00)	0.703 (0.04)	0.208 (0.02)	0.989 (0.00)	0.857 (0.01)
24h Linear	0.917 (0.00)	0.665 (0.05)	0.186 (0.02)	0.989 (0.00)	0.857 (0.01)
18h GPm52	0.921 (0.00)	0.788 (0.04)	0.191 (0.02)	0.994 (0.00)	0.846 (0.01)
18h GPm32	0.920 (0.00)	0.796 (0.04)	0.181 (0.02)	0.994 (0.00)	0.845 (0.01)
18h GP	0.921 (0.00)	0.789 (0.04)	0.186 (0.02)	0.994 (0.00)	0.844 (0.01)
18h Linear	0.914 (0.00)	0.714 (0.04)	0.145 (0.02)	0.991 (0.00)	0.831 (0.01)
12h GP	0.914 (0.00)	0.737 (0.05)	0.101 (0.01)	0.994 (0.00)	0.799 (0.01)
24h SVM	0.912 (0.00)	0.798 (0.04)	0.089 (0.01)	0.994 (0.00)	0.793 (0.01)
12h GPm52	0.914 (0.00)	0.734 (0.05)	0.102 (0.01)	0.993 (0.00)	0.791 (0.01)
12h GPm32	0.913 (0.00)	0.727 (0.05)	0.101 (0.01)	0.993 (0.00)	0.787 (0.01)
12h Linear	0.910 (0.00)	0.701 (0.05)	0.076 (0.01)	0.992 (0.00)	0.779 (0.01)
12h SVM	0.910 (0.00)	0.969 (0.02)	0.000 (0.00)	0.999 (0.00)	0.669 (0.02)
18h SVM	0.909 (0.00)	0.935 (0.03)	0.012 (0.01)	0.999 (0.00)	0.597 (0.02)

Figure 2 shows the effect of varying the length of the time series. Increasing time series length improves the prediction result up to 48h for GP and the linear model. There is no marked difference between the 48h and 72h predictions. SVM performance peaks at 32h.

Using only time series data AUC is initially low, surpassing that of the SNAP-II/SNAPPE-II combination with 36h and longer time series. The best AUC was achieved with 48h data. Using the full 72h time series results in




Figure 1: ROC curves for 12h and 72h classifiers. The ROCs of SNAP-II, SNAPPE-II, and the random classifier are also shown. Top row: all available features. Middle row: time series data + GA + BW. Bottom row: time series data only.

235 slightly lower AUC scores. The addition of GA and BW improves the result with short time series (12h and 18h), but has little effect with 24h and longer




Figure 2: AUC values for different time series lengths. Top left: all variables. Top right: time series data with GA and BW. Bottom: time series data only.

time series.

4. Discussion

In another study [30], birth weight alone was found to have AUC 0.74 and gestational age alone had AUC 0.71. Both were inferior to the Clinical Risk Index for Babies (CRIB) [31], which had AUC 0.82. CRIB-II [32] was also found to have inferior predictive power at AUC 0.69. A comparison study of CRIB-II and SNAPPE-II [21] found the two scores performing equally well, with CRIB-II AUC 0.913 (SE 0.014) and SNAPPE-II AUC 0.907 (SE 0.012), while another comparison study [33] found CRIB (AUC 0.90) and CRIB-II (AUC 0.91) superior to SNAPPE-II (AUC 0.84, which is close to our SNAPPE-II thresholding result).

While our best prediction results were achieved using all available variables, adding SNAP-II and SNAPPE-II to time series data with gestational age at birth and birth weight did not markedly improve the results. This is not surprising given that both scores are influenced by GA and BW to a great extent.

Both GP classification and the linear probit model gave practically identical results with time series of 36h and longer. With time series data only, GP and the linear probit model lose to SNAP-II/SNAPPE-II thresholding with 12h and 18h data, achieve roughly equal performance with 24h data, and beat them with 36h and longer time series, as does SVM. However, SVM performs significantly worse than SNAP-II/SNAPPE-II with 12h to 24h data.

In Figure 2 there is a slight drop in performance when using 72h data instead of 48h data. Although this may look surprising, it could be explained by the fact that the feature computations did not explicitly take the length of the time interval into account. It is thus possible that by, for example, limiting the feature computations to the end of the time series or by taking the length of the time interval otherwise into account could improve the predictions.

The highest sensitivity achieved was 0.475 for the linear probit model using all variables with 72h data. GP sensitivities varied from 0.283 to 0.453 using all variables, dropping down markedly (0.101...0.361) with only time series data. It is worth noting that low sensitivities of predictions do not necessarily mean that the clinical value of the predictions is low. From the clinical viewpoint, specificity is more important than sensitivity when predicting mortality. If the clinicians suspect that there is a high risk that the preterm infant will die, this can affect decisions to perform risky operations or start resource-intensive treatments. These kinds of decisions require careful consideration of the clinical situation and never rely on a single factor, such as predictive models. The goal is to have as high specificity as possible to avoid withholding treatment.

The prediction of in-hospital death in itself is not something that would be a major factor in how to treat the patient, but it can be useful in deciding whether to use some heavy means of care such as complex operations which themselves can be a risk to the patient. For that reason we have chosen to use data from

the early phase of the NICU stay. In the early stages the medical personnel
 280 have not yet been able to form a complete view of the patient's state.

5. Conclusions

Time series data from the initial hours of a preterm infant's intensive care
 unit stay can be used to improve the accuracy of existing methods for predicting
 in-hospital death. A Bayesian Gaussian process classifier can be used to create
 285 a predictive model. Combining features extracted from time series data with
 clinical scores calculated on arrival gives classification results in excess of clinical
 standards. Using only time series data gives results comparable with existing
 clinical standards, given a long enough time series.

As current NICU patient data systems already collect sensory data used
 290 in this paper, predictive modeling could be included in the care process to
 give physicians advance warning of increased risk of in-hospital death. The
 model already outperforms existing methods in our retrospective cohort and
 with further refinement could prove to be a valuable clinical tool.

Acknowledgement. The authors would like to thank Aki Vehtari for help
 295 in Gaussian process classifier design and implementation, and the Academy of
 Finland (projects 295505 and 266940) for financial support. We acknowledge
 the computational resources provided by the Aalto Science-IT project.

References

- [1] N. Byrnes, Can Technology Fix Medicine?, MIT Technology Re-
 300 view (Sep/Oct 2014).
- [2] D. A. Clifton, K. E. Niehaus, P. Charlton, G. W. Colopy, Health Informat-
 ics via Machine Learning for the Clinical Management of Patients:, IMIA
 Yearbook 10 (1) (2015) 38–43. doi:10.15265/IY-2015-014.
- [3] D. You, L. Hug, S. Ejdemyr, J. Beise, Levels and trends in child mortality.
 305 Report 2015. Estimates developed by the UN Inter-agency Group for Child

Mortality Estimation., Tech. rep., United Nations Inter-agency Group for Child Mortality Estimation (UN IGME), New York, NY (2015).

- [4] R. M. Patel, S. Kandefer, M. C. Walsh, E. F. Bell, W. A. Carlo, A. R. Laptook, P. J. Sánchez, S. Shankaran, K. P. Van Meurs, M. B. Ball, E. C. Hale, N. S. Newman, A. Das, R. D. Higgins, B. J. Stoll, Causes and Timing of Death in Extremely Premature Infants from 2000 through 2011, *New England Journal of Medicine* 372 (4) (2015) 331–340. doi:10.1056/NEJMoa1403489.
- [5] D. K. Richardson, J. E. Gray, M. C. McCormick, K. Workman, D. A. Goldmann, Score for Neonatal Acute Physiology: A physiologic severity index for neonatal intensive care, *Pediatrics* 91 (3) (1993) 617–623.
- [6] D. K. Richardson, J. D. Corcoran, G. J. Escobar, S. K. Lee, SNAP-II and SNAPPE-II: Simplified newborn illness severity and mortality risk scores, *The Journal of Pediatrics* 138 (1) (2001) 92–100. doi:10.1067/mpd.2001.109608.
- [7] C. E. Rasmussen, C. K. I. Williams, Gaussian Processes for Machine Learning, The MIT Press, 2006.
- [8] C. Cortes, V. Vapnik, Support-vector networks, *Machine learning* 20 (3) (1995) 273–297.
- [9] J. C. Platt, Probabilistic Outputs for Support Vector Machines and Comparisons to Regularized Likelihood Methods, in: *Advances in Large Margin Classifiers*, MIT Press, 1999, pp. 61–74.
- [10] M. Alvarez, D. Luengo, N. Lawrence, Latent force models, in: *Artificial Intelligence and Statistics*, 2009, pp. 9–16.
- [11] M. A. Alvarez, D. Luengo, N. D. Lawrence, Linear Latent Force Models Using Gaussian Processes, *IEEE Transactions on Pattern Analysis and Machine Intelligence* 35 (11) (2013) 2693–2705. doi:10.1109/TPAMI.2013.86.

- [12] J. Hartikainen, S. Särkkä, Sequential inference for latent force models, Proceedings of The 27th Conference on Uncertainty in Artificial Intelligence (UAI 2011).
- [13] J. Hartikainen, M. Seppänen, S. Särkkä, State-space inference for non-linear latent force models with application to satellite orbit prediction, in: Proceedings of the 29th International Conference on Machine Learning (ICML-12), Edinburgh, Scotland, UK, 2012, pp. 903–910.
- [14] S. Särkkä, A. Solin, J. Hartikainen, Spatiotemporal learning via infinite-dimensional Bayesian filtering and smoothing: A look at Gaussian process regression through Kalman filtering, *IEEE Signal Processing Magazine* 30 (4) (2013) 51–61.
- [15] S. Särkkä, *Bayesian Filtering and Smoothing*, Cambridge University Press, 2013.
- [16] G. W. Colopy, M. A. Pimentel, D. A. Clifton, S. J. Roberts, Bayesian Gaussian Processes for Identifying the Deteriorating Patient, in: Proceedings of the 38th Annual Conference of the IEEE Engineering in Medicine and Biology Society, IEEE, Orlando, FL, USA, 16-20 Aug 2016, pp. 5311–5314. doi:10.1109/EMBC.2016.7591926.
- [17] M. Ghassemi, M. A. Pimentel, T. Naumann, T. Brennan, D. A. Clifton, P. Szolovits, M. Feng, A multivariate timeseries modeling approach to severity of illness assessment and forecasting in ICU with sparse, heterogeneous clinical data, in: Proceedings of the 29th AAAI Conference on Artificial Intelligence, Vol. 1, NIH Public Access, Austin, TX, USA, 2015, pp. 446–453.
- [18] F. Güiza, J. Ramon, H. Blockeel, Gaussian processes for prediction in intensive care, in: *Gaussian Processes in Practice Workshop*, Bletchley Park, UK, 2006, pp. 1–4.

- [19] V. Gangadharan, Automated multi-parameter monitoring of neonates, PhD thesis, UCL (University College London), London (2013).
- [20] S. Saria, A. K. Rajani, J. Gould, D. Koller, A. A. Penn, Integration of Early Physiological Responses Predicts Later Illness Severity in Preterm Infants, *Science Translational Medicine* 2 (48) (2010) 48ra65–48ra65, doi: 10.1126/scitranslmed.3001304.
- [21] S. Reid, B. Bajuk, K. Lui, E. A. Sullivan, NSW and ACT Neonatal Intensive Care Units Audit Group, PSN, Comparing CRIB-II and SNAPPE-II as mortality predictors for very preterm infants: Comparing CRIB-II and SNAPPE-II, *Journal of Paediatrics and Child Health* 51 (5) (2015) 524–528. doi:10.1111/jpc.12742.
- [22] O.-P. Rinta-Koski, S. Särkkä, J. Hollmén, M. Leskinen, S. Andersson, Prediction of preterm infant mortality with Gaussian process classification, in: *Proceedings of the 25th European Symposium on Artificial Neural Networks, Computational Intelligence and Machine Learning*, Bruges, Belgium, 26-28 April 2017, pp. 193–198.
- [23] J. Vanhatalo, J. Riihimäki, J. Hartikainen, P. Jylänki, V. Tolvanen, A. Vehtari, GPstuff: Bayesian modeling with Gaussian processes, *Journal of Machine Learning Research* 14 (Apr) (2013) 1175–1179.
- [24] A. Solin, S. Särkkä, The 10th annual MLSP competition: First place, in: *2014 IEEE International Workshop on Machine Learning for Signal Processing (MLSP)*, IEEE, 2014, pp. 1–3.
- [25] K. Suotsalo, S. Särkkä, Detecting Malignant Ventricular Arrhythmias in Electrocardiograms by Gaussian Process Classification, in: *Proceedings of the 27th IEEE International Workshop on Machine Learning for Signal Processing (MLSP)*, IEEE, Tokyo, Japan, September 25-28, 2017.
- [26] J. A. Swets, Measuring the Accuracy of Diagnostic Systems, *Science* 240 (1988) 1285–1293.

- [27] A. P. Bradley, The use of the area under the ROC curve in the evaluation of machine learning algorithms, *Pattern recognition* 30 (7) (1997) 1145–1159.
- [28] W. J. Youden, Index for rating diagnostic tests, *Cancer* 3 (1950) 32–35. doi:10.1002/1097-0142(1950)3:1<32::AID-CNCR2820030106>3.0.CO;2-3.
- [29] H. Rue, S. Martino, N. Chopin, Approximate Bayesian inference for latent Gaussian models by using integrated nested Laplace approximations, *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 71 (2) (2009) 319–392. doi:10.1111/j.1467-9868.2008.00700.x.
- [30] C. Bühner, B. Metze, M. Obladen, CRIB, CRIB-II, birth weight or gestational age to assess mortality risk in very low birth weight infants?, *Acta Paediatrica* 97 (7) (2008) 899–903. doi:10.1111/j.1651-2227.2008.00793.x.
- [31] The International Neonatal Network, The CRIB (clinical risk index for babies) score: A tool for assessing initial neonatal risk and comparing performance of neonatal intensive care units, *The Lancet* 342 (8865) (1993) 193–198.
- [32] G. Parry, J. Tucker, W. Tarnow-Mordi, CRIB II: An update of the clinical risk index for babies score, *The Lancet* 361 (9371) (2003) 1789–1791.
- [33] L. Gagliardi, A. Cavazza, A. Brunelli, M. Battaglioli, D. Merazzi, F. Tandoi, D. Cella, G. F. Perotti, M. Pelti, I. Stucchi, F. Frisone, A. Avanzini, R. Bellù, the NNL study group, Assessing mortality risk in very low birth-weight infants: A comparison of CRIB, CRIB-II, and SNAPPE-II, *Archives of Disease in Childhood - Fetal and Neonatal Edition* 89 (5) (2004) F419–F422. doi:10.1136/adc.2003.031286.

Author's Biography



415

Olli-Pekka Rinta-Koski, Lic.Sc. (Tech.), is a PhD student at the Department of Computer Science, Aalto University School of Science. His research interests are the applications of machine learning in the fields of medicine, health, and well-being.



420

Simo Särkkä, D.Sc. (Tech.), is an Associate Professor and Academy Research Fellow with Aalto University, Technical Advisor of IndoorAtlas Ltd., and an Adjunct Professor with Tampere University of Technology and Lappeen-

ranta University of Technology. His research interests are in multi-sensor data
 425 processing systems with applications in location sensing, health technology,
 machine learning, inverse problems, and brain imaging.



Jaakko Hollmén, D.Sc. (Tech.), is a senior researcher at the Department
 of Computer Science, Aalto University School of Science and the group leader
 430 of the Parsimonious Modelling research group at Helsinki Institute for Infor-
 mation Technology HIIT. The research group develops computational meth-
 ods for data analysis with applications in cancer genomics and environmental
 informatics. His research interests include machine learning, data mining, and
 their applications in bioinformatics and environmental time series analysis.



Markus Leskinen, MD, PhD, is a neonatologist at Children's Hospital,
 University of Helsinki, and Helsinki University Hospital, Helsinki, Finland. He

studied medicine at the University of Helsinki and specialized in pediatrics. His
re- search interests include big data in neonatology and advances in quality of
440 care.



Sture Andersson, MD, PhD, is a pediatrician and neona- tologist by train-
ing. He is a professor of neonatology at Chil- drens Hospital, University of
Helsinki, and Helsinki Univer- sity Hospital, Helsinki, Finland. His main re-
445 search areas are lung injury in preterm infants, postnatal pulmonary adapta-
tion, and big data in neonatology